


Optimal ε^* vs a poorly-chosen $\varepsilon = 0.05$



SKR improvement from optimising ε
(mild baseline $\varepsilon = 0.01$)

